

DeepSpeed

EEE549: Final Presentation

Kaiyuan Tan, Perfect Obumneme, Emmanuel Adelaju, Sung Park, Atharva Hundare

School of Electrical, Computer, and Energy Engineering,
Arizona State University

Contacts: [ktan24, pobumnem, eadelaju, spark259, ahundare]@asu.edu

Previous Reward Function (First Contest)

① Strict Off-Track Penalty

if off-track $\Rightarrow R = 0$

② Lap Completion Bonus

if progress = 100% $\Rightarrow R \leftarrow R + 100$

③ Cubic Speed Factor

Let Max speed = v_p , Current Speed = v

$$g(v) = \begin{cases} 0.5, & v < 1 \\ (\max(0.1, \min(1, v/v_p)))^3, & v \geq 1 \end{cases}$$

$$R \leftarrow R \cdot g(v)$$

④ On-track Constraints

d_c = Distance from center, W = Track Width,

$$R \leftarrow \begin{cases} R + 1, & d_c \leq 0.4W \\ 0.3R, & d_c > 0.4W \end{cases}$$


⑤ Pose Control (Future Direction)

$$\Delta\theta = \text{wrap}_{[-180, 180]}(\theta_{\text{track}} - \theta_h)$$

If $|\Delta\theta| > 10^\circ$:

$$e = \frac{|\Delta\theta|}{15^\circ}, \quad m = \max(0.01, 1 - e^2)$$

$$R \leftarrow R \cdot m$$



⑥ Sharp Steering Penalty

if $|\text{steering}| > 20^\circ$ and $v > 2.5 \Rightarrow R \leftarrow 0.7R$

First Contest Rank

Rank	Racer	Qualifying time	Gap to 1st	Video	Off-track
1	AccessibleInterface#1358	00:39.199	-	View	10
2	Fast--and--Curious	00:49.395	+00:10.196	View	12
3	TheExplorers	00:50.462	+00:11.263	View	20
4	Horse	00:53.999	+00:14.800	View	8
5	SupervisedBySchumacher	00:55.865	+00:16.666	View	20
6	The-Reinforcers	00:56.927	+00:17.728	View	6
7	PathingNumber#6148	00:57.534	+00:18.335	View	6
8	ModelForgers	00:58.134	+00:18.935	View	14
9	DeepSpeed21	01:03.136	+00:23.937	View	16
10	LoopingCoder#4099	01:08.468	+00:29.269	View	11
11	SequentialMachine#5150	01:15.797	+00:36.598	View	21
12	AbstractRequirements#3057	01:21.463	+00:42.264	View	11
13	MemoryResource#0626	01:21.672	+00:42.473	View	17

Final Reward Function



Foundational Geometry



Track Data Preparation



Target Selection:
Lookahead Strategy



Deriving the Optimal
Steering Angle



Reward Scoring
and Evaluation



**Our Final Reward function
Sub-components**

Foundational Geometry

Our agent is solving a **geometry problem** every timestep.

We represent the car as a point in \mathbb{R}^2 with a heading:

$$C = (x_c, y_c, \psi)$$


The track is a sequence of waypoints:

$$W = \{(x_i, y_i)\}$$

Everything our policy learns flow from 3 primitives


① Distance

$$\sqrt{(x_2 - x_1)^2 + (y_2 - y_1)^2}$$



② Direction (Polar)

$$(r, \theta) = \left(\sqrt{x^2 + y^2}, \arctan 2(y, x) \right)$$



③ Shortest Rotation

$$\Delta\theta = ((\Delta\theta_{raw} + 180^\circ) \bmod 360^\circ) - 180^\circ$$

Car heading: $\psi = 10^\circ$, Target direction: $\theta = 300^\circ$

Raw angle difference: $\Delta\theta_{raw} = \theta - \psi = 300^\circ - 10^\circ = 290^\circ$

This suggests turning $+290^\circ$, which is obviously wrong!
That's a huge spin.

Apply normalization:

$$\text{Result: } \Delta\theta = 110^\circ - 180^\circ = -70^\circ$$

← Turn left 70° , not right 290°

Track Data Preparation

Deepracer tracks **sparse waypoints** in a circuit.

Between two waypoints, the true path might **curve**, **dip inward**, or **swing wide** – but the agent only see **two straight-line dots**.

Upsampling waypoints makes the track feel **continuous**:

We took each segment: $P_i = (x_i, y_i)$, $P_{i+1} = (x_{i+1}, y_{i+1})$

Then generated intermediate points using linear interpolation:

$$P_{i,k} = (1 - \lambda)P_i + \lambda P_{i+1}, \quad \lambda = \frac{k}{N}$$

Where:


$N = 20$ (upsample factor)

$k \in \{0, 1, 2, \dots, N - 1\}$


This yields **20 mini-waypoints per original segment**.

In code terms: *densify the map by a factor of 20*.

Sparse
Waypoints



Upsampled
Track



2022 Summit Speedway

Target Selection: Lookahead Strategy

- 1 Firstly, we collect the upsampled waypoints and the position of the car

$$C = (x_c, y_c)$$

Car location

$$W = \{w_0, w_1, w_2, \dots\}$$

We reorder the waypoint list so that waypoint 0 is the closest to the car using our primitives

- 2 We define a lookahead radius

Radius, $r = 0.9 \cdot \text{track_width}$

Conceptually: a circle centered at the car.


Every waypoint inside the circle is “too immediate or noisy”.

We want a point just beyond this bubble to aim at.

- 3 Drive forward and find target

For each waypoint w_i in the driving order:

$$d(w_i, C) = \sqrt{(x_i - x_c)^2 + (y_i - y_c)^2}$$




Stop at the first waypoint where: $d(w_i, C) > r$

$\text{Target} = w_j \text{ s.t. } d(w_j, C) > r \text{ and } j \text{ is minimal.}$

All edge cases fallback to the closest waypoint.

Deriving the Optimal Steering Angle

```
def get_target_steering_degree(params):
    tx, ty = get_target_point(params)
    car_x = params['x']
    car_y = params['y']
    dx = tx - car_x
    dy = ty - car_y
    heading = params['heading']
    _, target_angle = polar(dx, dy)
    steering_angle = target_angle - heading
    return angle_mod_360(steering_angle)
```



We chose our target, $T = (t_x, t_y)$

Determine car position, $C = (x_c, y_c, \psi)$

Compute the target-relative vector
 $\hat{v} = (dx, dy) = (t_x - x_c, t_y - y_c)$

Compare target direction vs car heading

We use our primitive to convert to polar coordinates:
 $(r, \theta) = polar(dx, dy) = (\sqrt{x^2 + y^2}, \arctan 2(y, x))$

Car heading: $\psi = params[heading]$,
Raw steering request: $\Delta\theta_{raw} = \theta - \psi$

Normalize to shortest rotation (primitive):
s.t. $\Delta\theta \in [-180^\circ, 180^\circ]$

Reward Scoring and Evaluation

```
def score_steer_to_point_ahead(params):
    best_stearing_angle = get_target_steering_degree(params)
    steering_angle = params['steering_angle']
    error = (steering_angle - best_stearing_angle) / 60.0
    score = 1.0 - abs(error)
    return max(score, 0.01)
```

```
def reward_function(params):
    return float(score_steer_to_point_ahead(params))
```

$$Score = 1 - |e|$$

$$\begin{cases} |e| = 0 \Rightarrow \text{score} = 1.0 \text{ (perfect)} \\ |e| = 0.2 \Rightarrow \text{score} = 0.8 \text{ (pretty good)} \\ |e| = 0.6 \Rightarrow \text{score} = 0.4 \text{ (poor)} \\ |e| \geq 1.0 \Rightarrow \text{score} \leq 0 \text{ (bad)} \end{cases}$$

Final reward is $\max(\text{score}, 0.01)$ to avoid reward collapse

From our geometry pipeline:
we already know the best steering angle, θ_{best}

The simulator gives us the actual steering: θ_{car}

Compute the steering error: $e = \Delta\theta_{car} - \Delta\theta_{best}$

$$\begin{cases} > 0 \rightarrow \text{car is turning too much to the right} \\ < 0 \rightarrow \text{car is turning too much to the left} \\ = 0 \rightarrow \text{perfect alignment} \end{cases}$$

We normalize the error: $\frac{e}{60}$

$$\begin{cases} \pm 60^\circ \text{ mismatch} = \text{total failure} \\ < \pm 10^\circ \text{ mismatch} = \text{good} \\ < \pm 5^\circ \text{ mismatch} = \text{excellent} \end{cases}$$

The car is rewarded for being close to ideal, not simply surviving.

Reward Function Highlight

- We do **NOT** incentive our Agent for **SPEED**
- We do **NOT** incentive our Agent for **wheels ONTRACK**
- We do **NOT** incentive our Agent for **PROGRESS**
- We **ONLY** incentive our Agent for **good STEERING decisions**

Observations:

A car that aligns its heading with the future of the track will do this:

- Enter corners earlier
- Cut apexes naturally
- Exit with straightened steering


These are the recipe for **SPEED** and **STABILITY...**

Agent Configuration and Hyperparameters


Category	Item	Value
Action space	Type	Continuous
	Speed range	0.5m/s : 3.7 m/s
	Steering angle range	28° : – 27°
Framework	Library	TensorFlow
Reinforcement learning	Algorithm	PPO

Hyperparameter	Value
Gradient descent batch size	64
Entropy	0.01
Discount factor	0.99 → 0.7
Loss type	Huber
Learning rate	0.0003
Number of experience episodes between each policy-updating iteration	20
Number of Epochs	10

Training Results









Training Track
Circuit de Barcelona-Catalunya




Training Time
190 minutes

Evaluation (Stable on Different Tracks)




Evaluation (Our Lap Time vs. WR Lap Time by Track)



- **Best Lap Time:**
Shanghai Sudu Training
 - ~7% slower
- **Worst Lap Time:**
Kuei Raceway
 - ~32% slower

Final Contest Performance (Offtrack vs. Lap Time)




This plot shows the tradeoff
Between raw speed and off-
track penalties.

Teams that drove cautiously
sacrificed lap time.

Teams that pushed too hard
accumulated resets.

Top performers found the balance -
Need for speed is the extreme outlier.

Future Reward Function Iteration



Corner Stability:

Reduce over-rotations and rare spin-outs on high curvature by penalizing sudden steering jumps

Smooth Acceleration:

Incentivize controlled throttle change to prevent slip

Speed Awareness:

Introduce reward for speed but only when aligned to the correct steering angle

Lap Completion Bonus:

Reward fast laps only after finishing the full

Team Member Contributions

Team Member	Contribution
Kaiyuan Tan	<ul style="list-style-type: none">• Create and train models with different reward functions
Perfect Obumneme	<ul style="list-style-type: none">• Iterated through several models and reward functions• Typeset all PPT equations and plotted PPT graphs
Emmanuel Adelaju	<ul style="list-style-type: none">• Ran different models on different tracks using different reward functions
Sung Park	<ul style="list-style-type: none">• Tried different approaches for reward functions and finetuning.• Worked on some Evaluation, and Car-model and Hyperparameters PPT
Atharva Hundare	<ul style="list-style-type: none">• Performed the final reward function design, training and evaluation on various tracks.

THANK YOU!